

Stock Market Prediction And Forecasting Using Hybrid QML Model(Quantum-Machine Learning!)

```
In [1]: ### Keras and Tensorflow >2.0
```

```
In [2]: ### Data Collection
import pandas_datareader as pdr
import yfinance as yf

df = yf.Ticker("AAPL")

# get historical market data
df = df.history(period="max")
```

```
In [3]: df
```

Out[3]:

	Open	High	Low	Close	Volume	Dividends	Stock Splits
Date							
1980-12-12	0.100453	0.100890	0.100453	0.100453	469033600	0.0	0.0
1980-12-15	0.095649	0.095649	0.095213	0.095213	175884800	0.0	0.0
1980-12-16	0.088661	0.088661	0.088224	0.088224	105728000	0.0	0.0
1980-12-17	0.090408	0.090845	0.090408	0.090408	86441600	0.0	0.0
1980-12-18	0.093029	0.093466	0.093029	0.093029	73449600	0.0	0.0
...
2021-11-15	150.369995	151.880005	149.429993	150.000000	59222800	0.0	0.0
2021-11-16	149.940002	151.490005	149.339996	151.000000	59256200	0.0	0.0
2021-11-17	151.000000	155.000000	150.990005	153.490005	88807000	0.0	0.0
2021-11-18	153.710007	158.669998	153.050003	157.869995	137827700	0.0	0.0
2021-11-19	157.649994	161.020004	156.529999	160.550003	117147500	0.0	0.0

10324 rows × 7 columns

```
In [4]: df.to_csv('AAPL.csv')
```

```
In [5]: import pandas as pd
```

```
In [6]: df=pd.read_csv('AAPL.csv')
```

```
In [7]: df.head()
```

Out[7]:

	Date	Open	High	Low	Close	Volume	Dividends	Stock Splits
0	1980-12-12	0.100453	0.100890	0.100453	0.100453	469033600	0.0	0.0
1	1980-12-15	0.095649	0.095649	0.095213	0.095213	175884800	0.0	0.0
2	1980-12-16	0.088661	0.088661	0.088224	0.088224	105728000	0.0	0.0
3	1980-12-17	0.090408	0.090845	0.090408	0.090408	86441600	0.0	0.0
4	1980-12-18	0.093029	0.093466	0.093029	0.093029	73449600	0.0	0.0

```
In [8]: df.tail()
```


Out[8]:

	Date	Open	High	Low	Close	Volume	Dividends	Stock Splits
10319	2021-11-15	150.369995	151.880005	149.429993	150.000000	59222800	0.0	0.0
10320	2021-11-16	149.940002	151.490005	149.339996	151.000000	59256200	0.0	0.0
10321	2021-11-17	151.000000	155.000000	150.990005	153.490005	88807000	0.0	0.0
10322	2021-11-18	153.710007	158.669998	153.050003	157.869995	137827700	0.0	0.0
10323	2021-11-19	157.649994	161.020004	156.529999	160.550003	117147500	0.0	0.0

```
In [9]: df1=df.reset_index()['Close']
```

```
In [10]: import matplotlib.pyplot as plt
plt.plot(df1)
```

```
Out[10]: [<matplotlib.lines.Line2D at 0x7ffac802f1f0>]
```



```
In [11]: ### LSTM are sensitive to the scale of the data. so we apply MinMax scaler
```

```
In [12]: import numpy as np
```

```
In [13]: df1
```

```
Out[13]: 0      0.100453
1      0.095213
2      0.088224
3      0.090408
4      0.093029
...
10319  150.000000
10320  151.000000
10321  153.490005
10322  157.869995
10323  160.550003
Name: Close, Length: 10324, dtype: float64
```

```
In [14]: from sklearn.preprocessing import MinMaxScaler
scaler=MinMaxScaler(feature_range=(0,1))
df1=scaler.fit_transform(np.array(df1).reshape(-1,1))
```

```
In [15]: print(df1)
```

```
[[3.86384327e-04]
 [3.53734058e-04]
 [3.10195539e-04]
 ...
 [9.56015647e-01]
 [9.83303335e-01]
 [1.00000000e+00]]
```

```
In [16]: ##splitting dataset into train and test split
training_size=int(len(df1)*0.65)
test_size=len(df1)-training_size
train_data,test_data=df1[0:training_size:],df1[training_size:len(df1),:1]
```

```
In [17]: training_size,test_size
```

```
Out[17]: (6710, 3614)
```

```
In [18]: train_data
```

```
Out[18]: array([[0.00038638],
 [0.00035373],
 [0.0003102 ],
 ...,
 [0.02603478],
 [0.02610536],
 [0.02625989]])
```

```
In [19]: import numpy

# This is like feature matrix.
def create_dataset(dataset, time_step=1):
```

```

dataX, dataY = [], []
for i in range(len(dataset)-time_step-1):
    a = dataset[i:(i+time_step), 0]    ###i=0, 0,1,2,3-----99    100
    dataX.append(a)
    dataY.append(dataset[i + time_step, 0])
return numpy.array(dataX), numpy.array(dataY)

```

```

In [20]: # reshape into X=t,t+1,t+2,t+3 and Y=t+4
time_step = 100 # this is like X = [t,t+1,t+2,t+3,t+4,t+5] and Y=[t+6]
X_train, y_train = create_dataset(train_data, time_step)
X_test, ytest = create_dataset(test_data, time_step)

```

```

In [21]: print(X_train.shape), print(y_train.shape)

```

```

(6609, 100)
(6609,)
(None, None)

```

```

In [22]: X_test.shape,ytest.shape

```

```

Out[22]: ((3513, 100), (3513,))

```

```

In [23]: # OK so this reshape is required, because we are sending it to LSTM! This is just reshape for now! :D

```

```

# reshape input to be [samples, time steps, features] which is required for LSTM
X_train =X_train.reshape(X_train.shape[0],X_train.shape[1] , 1)
X_test = X_test.reshape(X_test.shape[0],X_test.shape[1] , 1)
X_train.shape,X_test.shape

# So here we have 6609 recorda(Rows) and 100 Columns(aka feature-Vectors) What is 1?

```

```

Out[23]: ((6609, 100, 1), (3513, 100, 1))

```

```

In [24]: ### Create the Stacked LSTM model
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from tensorflow.keras.layers import LSTM

```

```

In [25]: model=Sequential()
model.add(LSTM(50,return_sequences=True,input_shape=(100,1)))
model.add(LSTM(50,return_sequences=True))
model.add(LSTM(50))
model.add(Dense(4))
model.add(Dense(1))
model.compile(loss='mean_squared_error',optimizer='adam')

```

```

2021-11-21 16:09:21.757088: I tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:939] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero
2021-11-21 16:09:21.762396: I tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:939] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero
2021-11-21 16:09:21.762950: I tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:939] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero
2021-11-21 16:09:21.763723: I tensorflow/core/platform/cpu_feature_guard.cc:151] This TensorFlow binary is optimized with oneAPI Deep Neural Network Library (oneDNN) to use the following CPU instructions in performance-critical operations:  AVX2 FMA
To enable them in other operations, rebuild TensorFlow with the appropriate compiler flags.
2021-11-21 16:09:21.765135: I tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:939] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero
2021-11-21 16:09:21.765572: I tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:939] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero
2021-11-21 16:09:21.765900: I tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:939] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero
2021-11-21 16:09:22.150949: I tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:939] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero
2021-11-21 16:09:22.151242: I tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:939] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero
2021-11-21 16:09:22.151493: I tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:939] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero
2021-11-21 16:09:22.151729: I tensorflow/core/common_runtime/gpu/gpu_device.cc:1525] Created device /job:localhost/replica:0/task:0/device:GPU:0 with 3493 MB memory:  -> device: 0, name: NVIDIA GeForce GTX 1660 Ti with Max-Q Design, pci bus id: 0000:01:00:0, compute capability: 7.5

```

```

In [26]: model.summary()

```

```

Model: "sequential"

```

Layer (type)	Output Shape	Param #
=====		

lstm (LSTM)	(None, 100, 50)	10400
lstm_1 (LSTM)	(None, 100, 50)	20200
lstm_2 (LSTM)	(None, 50)	20200
dense (Dense)	(None, 4)	204
dense_1 (Dense)	(None, 1)	5

```

=====
Total params: 51,009
Trainable params: 51,009
Non-trainable params: 0

```

```
In [27]: model.fit(X_train,y_train,validation_data=(X_test,ytest),epochs=10,batch_size=64,verbose=1)
```

Epoch 1/10

2021-11-21 16:09:26.019054: I tensorflow/stream_executor/cuda/cuda_dnn.cc:366] Loaded cuDNN version 8300

104/104 [=====] - 7s 32ms/step - loss: 3.9244e-05 - val_loss: 0.0190

Epoch 2/10

104/104 [=====] - 3s 26ms/step - loss: 2.8158e-07 - val_loss: 0.0033

Epoch 3/10

104/104 [=====] - 3s 25ms/step - loss: 1.5459e-07 - val_loss: 0.0031

Epoch 4/10

104/104 [=====] - 3s 25ms/step - loss: 1.9859e-07 - val_loss: 0.0029

Epoch 5/10

104/104 [=====] - 3s 25ms/step - loss: 2.0481e-07 - val_loss: 0.0030

Epoch 6/10

104/104 [=====] - 3s 25ms/step - loss: 1.6410e-07 - val_loss: 0.0028

Epoch 7/10

104/104 [=====] - 3s 25ms/step - loss: 2.0121e-07 - val_loss: 0.0023

Epoch 8/10

104/104 [=====] - 3s 26ms/step - loss: 1.9869e-07 - val_loss: 0.0038

Epoch 9/10

104/104 [=====] - 3s 26ms/step - loss: 1.7274e-07 - val_loss: 0.0026

Epoch 10/10

104/104 [=====] - 3s 25ms/step - loss: 1.5691e-07 - val_loss: 0.0025

```
Out[27]: <keras.callbacks.History at 0x7ffa58252bb0>
```

```
In [28]: ### Lets Do the prediction and check performance metrics
train_predict=model.predict(X_train)
test_predict=model.predict(X_test)
```

```
In [29]: ##Transformback to original form
train_predict=scaler.inverse_transform(train_predict)
test_predict=scaler.inverse_transform(test_predict)
```

```
In [30]: ### Calculate RMSE performance metrics
import math
from sklearn.metrics import mean_squared_error
math.sqrt(mean_squared_error(y_train,train_predict))
```

```
Out[30]: 0.7679865322838249
```

```
In [31]: ### Test Data RMSE
math.sqrt(mean_squared_error(ytest,test_predict))
```

```
Out[31]: 43.025365207938115
```

```
In [32]: import pennylane as qml

n_qubits = 2
dev = qml.device("default.qubit", wires=n_qubits)

@qml.qnode(dev)
def qnode(inputs, weights):
    qml.AngleEmbedding(inputs, wires=range(n_qubits))
    qml.BasicEntanglerLayers(weights, wires=range(n_qubits))
    return [qml.expval(qml.PauliZ(wires=i)) for i in range(n_qubits)]

n_layers = 6
weight_shapes = {"weights": (n_layers, n_qubits)}

qlayer = qml.qnn.KerasLayer(qnode, weight_shapes, output_dim=n_qubits)

model=Sequential()
```

```

model.add(LSTM(50,return_sequences=True,input_shape=(100,1)))
model.add(LSTM(20,return_sequences=True))
model.add(LSTM(10))
model.add(Dense(2))
model.add(qlayer)
model.add(Dense(1))
model.compile(loss='mean_squared_error',optimizer='adam')
model.summary()

```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
lstm_3 (LSTM)	(None, 100, 50)	10400
lstm_4 (LSTM)	(None, 100, 20)	5680
lstm_5 (LSTM)	(None, 10)	1240
dense_2 (Dense)	(None, 2)	22
keras_layer (KerasLayer)	(None, 2)	0 (unused)
dense_3 (Dense)	(None, 1)	3

```

=====
Total params: 17,345
Trainable params: 17,345
Non-trainable params: 0

```

In [33]: `model.fit(X_train,y_train,validation_data=(X_test,ytest),epochs=2,batch_size=64,verbose=1)`

```

Epoch 1/2
104/104 [=====] - 570s 5s/step - loss: 0.0013 - val_loss: 0.0990
Epoch 2/2
104/104 [=====] - 610s 6s/step - loss: 1.3523e-05 - val_loss: 0.0893
<keras.callbacks.History at 0x7ff9e00925b0>

```

Out[33]:

In [37]: `model.summary()`

Model: "sequential_1"

Layer (type)	Output Shape	Param #
lstm_3 (LSTM)	(None, 100, 50)	10400
lstm_4 (LSTM)	(None, 100, 20)	5680
lstm_5 (LSTM)	(None, 10)	1240
dense_2 (Dense)	(None, 2)	22
keras_layer (KerasLayer)	(None, 2)	12
dense_3 (Dense)	(None, 1)	3

```

=====
Total params: 17,357
Trainable params: 17,357
Non-trainable params: 0

```

In [34]: `### Lets Do the prediction and check performance metrics`
`train_predict=model.predict(X_train)`
`test_predict=model.predict(X_test)`

In [35]: `##Transformback to original form`
`train_predict=scaler.inverse_transform(train_predict)`
`test_predict=scaler.inverse_transform(test_predict)`

`### Calculate RMSE performance metrics`
`import math`
`from sklearn.metrics import mean_squared_error`
`math.sqrt(mean_squared_error(y_train,train_predict))`

Out[35]: 0.48442422883130176

In [36]: `# Calculate RMSE performance on Test Metrics!`
`math.sqrt(mean_squared_error(ytest,test_predict))`

Out[36]: 1.7449557868671541

